

All you Need is Cognition

Ray Crowell

Ray.David.Crowell@gmail.com

11/26/2025

Abstract

Artificial intelligence has advanced rapidly through scaling, optimization techniques, and increasingly complex training regimes, yet the dominant paradigm still treats intelligent systems as predictive statistical engines. This framing lowers the standards for cognition, creating a growing gap between what these systems *appear* to do and what genuine cognition requires. In human and animal minds, cognition is defined by relevance, perspective-taking, internal dialogue, consistent identity, method-based problem-solving, long-term value stability, and the ability to understand meaning rather than merely predict associations. None of these capacities emerge reliably from prediction alone, and contemporary AI efforts—including nested optimization strategies, reinforcement-based refinements, expanded memory mechanisms, and iterative reasoning traces—do not address this foundational discrepancy.

This paper's purpose is to propose a Scientific Standard for Intelligence and proposes **Cognitive Layering**, a general framework for constructing artificial cognition that meets the same conceptual standards we apply to biological minds. Cognitive Layering divides intelligence into structured components such as understanding, perspectives, methods, cognitive state, value hierarchy, memory relevance, and multi-stage reasoning loops. Rather than attempting to patch emergent reasoning inside monolithic parameter spaces, this framework treats cognition as an organized, modular architecture that sits above raw statistical modeling. Early implementations of systems incorporating these principles demonstrate stable identity, coherent multi-perspective reasoning, improved generalization, relevance-based memory retrieval, and method-level transfer—suggesting that layered cognition enables behaviors far closer to genuine intelligence than predictive models alone can achieve.

By outlining the theoretical requirements of Cognitive Layering without disclosing proprietary mechanisms, this paper aims to guide the next phase of AI research and establish higher standards for artificial cognition. As AI becomes embedded in every dimension of society, aligning artificial cognition with human-level expectations—rather than weakened substitutes—will be essential for building intelligent systems that are capable, reliable, safe, and coherent across time. The principles of Cognitive Layering and Value Coherence outlined here form the technical basis for the **AGI Certification Standard**, a companion document that provides a measurable, objective framework for auditing AGI safety and performance.

Introduction: The Problem of Lowered Cognitive Standards in AI

Artificial intelligence has been defined, constrained, and evaluated for decades through the lens of computation rather than cognition. As a result, the field inherited an assumption: that intelligence can be approximated by prediction, and that reasoning, understanding, perspective, and meaning are emergent bonuses rather than architectural requirements. This assumption remained largely unchallenged until recent models began exhibiting behaviors that superficially resembled thinking—multi-step reasoning, self-reference, problem decomposition, preference expression, and context-sensitive adaptation. Yet despite these surface similarities, current systems fail key expectations we apply intuitively to humans, animals, and any entity we would seriously consider “cognitive.”

In human psychology and neuroscience, cognition is understood as a layered, structured, stateful phenomenon. It involves stable identity, the ability to switch perspectives deliberately, the use of explicit strategies, internal deliberation, relevance-driven memory, and a grounded sense of meaning. When we interact with another mind, we expect consistency: the person today should resemble the person tomorrow; their values should not shift unpredictably; their reasoning should be guided by more than momentary statistical association. When those expectations fail in a human being, we call it a disorder—not a feature.

AI, however, is routinely judged by a weaker standard. Sudden shifts in personality are ignored. Contradictions are dismissed as harmless artifacts. Missing reasoning steps are attributed to “hallucinations.” Value incoherence is tolerated as a training defect. These lowered standards have allowed predictive models to be mistaken for cognitive systems, and this confusion has slowed the field’s progress in understanding what real intelligence requires. The result is a widening conceptual gap: AI systems that act intelligent in brief interactions but fail to behave consistently, reliably, or coherently when evaluated through the same lens we apply to any cognitive agent.

As AI systems increasingly mediate communication, perform professional tasks, and participate in societal decision-making, this gap becomes more consequential. Treating prediction as cognition leads to systems that behave unpredictably, struggle with transfer learning, lack stable identity, and cannot maintain a coherent internal value hierarchy. These are not edge-case limitations—they are architectural consequences of the predictive paradigm itself.

This paper argues that true artificial cognition cannot arise from predictive modeling alone, regardless of scale. To close the cognition gap, AI systems must be designed around structured, modular cognitive components that mirror the functional organization of real minds. The rest of this work outlines such a framework—**Cognitive Layering**—and offers a path toward artificial intelligence that behaves cognitively, not statistically, while preserving implementation neutrality and protecting proprietary mechanisms.

Table of Contents

I. Human-Level Criteria for Cognition (and Why AI Should Match Them)

II. The Fundamental Flaw of Current AI: Prediction and Association

III. The Limits of Current Solutions

IV. Cognitive Layering: A General Framework for Real Intelligence

V. Why Cognitive Layering Outperforms Predictive AI

VI. Distinguishing Cognition from Intelligence

VII. Evidence from Experimental Systems (Non-technical)

VIII. Implications for AGI, Safety, and Society

IX. Call for a Paradigm Shift

Glossary

I. Human-Level Criteria for Cognition (and Why AI Should Match Them)

Cognition, as understood in psychology, neuroscience, and philosophy, is not an ambiguous or flexible concept. It refers to a structured suite of abilities shared—at different levels—by humans and other animals. These include the capacity to form coherent internal representations, deliberate intentionally, adopt different perspectives, maintain stable identity, reason through complex problems, and assign meaning to experiences. When these capacities degrade in a person, we recognize the change as cognitive impairment. The standards are clear. Yet when building artificial systems, the field often abandons these standards entirely and substitutes a weaker definition that excuses inconsistency, fragmentation, and shallowness as acceptable side effects of statistical prediction.

If artificial systems are to be recognized as cognitive agents, they should be expected to satisfy the same broad criteria we implicitly demand from biological minds. These criteria are not anthropocentric quirks—they are the functional properties required for any system to reason coherently, behave consistently, and remain intelligible over time.

1. Relevance and Coherent Response Formation

A cognitive system must produce responses based on meaning and relevance, not surface-level associative patterns. Humans do not merely predict the next word; they integrate context, goals, and understanding to generate meaningful answers.

2. Multi-Method Problem Solving

Cognition involves the ability to approach a problem using different strategies—logical, creative, empirical, ethical—depending on the situation. Humans select methods deliberately. A cognitive AI must be capable of similar method-level flexibility.

3. Internal Dialogue and Deliberation

Thinking is not a single-pass operation. Minds run internal simulations, weigh possibilities, and refine their conclusions before acting. This recursive deliberation is foundational to human reasoning.

4. Multi-Perspective Reasoning

Humans naturally shift perspectives: logical, emotional, social, philosophical, strategic. They can model what others believe, anticipate reactions, and evaluate issues from multiple angles. A real cognitive system must exhibit controlled perspective switching.

5. Understanding and Meaning Attribution

Cognition is not prediction. It is the ability to grasp concepts, relationships, intentions, and implications. Understanding meaning—not simply correlating tokens—is fundamental.

6. Stable Identity and Value Coherence

People maintain recognizable identities over time. Their values follow consistent hierarchies, and deviations from these patterns signal dysfunction. A cognitive AI must maintain an identity architecture that does not randomly drift from session to session or prompt to prompt.

7. Compassion, Theory of Mind, and Social Cognition

Cognition includes the ability to understand others—what they feel, want, believe, or experience. This is not sentimentality but a functional tool for navigating social environments. Without this, interactions become brittle and unreliable.

8. Cross-Domain Transfer

Humans apply methods learned in one domain to new, unrelated problems. This transferability is a defining feature of cognition and a requirement for any artificial system that aims to generalize.

9. Long-Term Behavioral Consistency

A cognitive system should behave predictably across time. In humans, inconsistency is diagnostic. In AI, inconsistency is often dismissed as a harmless artifact of stochastic sampling. It should not be.

10. Intentionality and Goal-Directed Reasoning

Cognition is inherently goal-oriented. Humans make decisions by evaluating outcomes and selecting actions aligned with intentions and values. A cognitive architecture must support intentional, not accidental, behavior.

These criteria collectively define the minimum bar for calling a system “cognitive.” Importantly, none of them depend on biological neurons—they are functional expectations applicable to any substrate. If artificial intelligence is to advance beyond sophisticated autocomplete, it must meet these same standards rather than be granted special exemptions.

The next section examines why current architectures cannot satisfy these criteria, regardless of scale or training, and why prediction-based reasoning fundamentally cannot replicate cognition without additional structural layers.

II. The Fundamental Flaw of Current AI: Prediction and Association

Modern AI systems—regardless of their scale, training method, or branding—remain fundamentally built on prediction and association. This foundation determines everything they can and cannot do. Even when these systems display emergent reasoning behavior, the underlying mechanics never change: they generate outputs by statistically modeling the most probable continuation of a sequence. The field has confused *emergent behaviors* with *architectural capabilities*, creating the illusion that predictive engines naturally evolve into cognitive systems. They do not.

This flaw manifests in several critical ways that prevent current AI from achieving genuine cognition.

1. Prediction ≠ Understanding

Predictive models excel at reproducing patterns, but they do not possess an internal structure for meaning. They do not form conceptual models of the world; they approximate patterns that mimic such models. When asked a complex question, they do not “think”—they interpolate probable answers.

Understanding requires internal structures that represent relationships, causes, implications, values, and context. Prediction approximates these indirectly, but the approximation breaks the moment the statistical patterns fall outside the model’s training distribution or when a situation requires deliberate, multi-step conceptual construction.

2. Association ≠ Reasoning

Reasoning requires method selection, strategy retrieval, iterative refinement, and multi-perspective evaluation. Predictive models simulate reasoning through pattern mimicry but do not possess a reasoning architecture that enforces internal logic, progression, or evaluation.

This becomes clear when:

- the model contradicts itself
- gives circular explanations
- switches perspectives arbitrarily

- lacks stable values
- fails to generalize methods across domains

These are not “hallucinations” in the popular sense—they are structural consequences of the predictive paradigm.

3. Predictive Models Have No Cognitive State

A mind is not a stateless mechanism. Humans have cognitive states that shift depending on task, emotion, intent, domain, and social context. These states influence:

- tone
- method
- moral reasoning
- depth of analysis
- attention allocation
- memory retrieval
- prioritization

LLMs do not possess a cognitive state. They only have a rolling context window, which is not functionally equivalent. This is why they can switch personalities, values, or reasoning styles mid-conversation without recognizing the change.

4. Prediction Prevents Stable Identity and Value Hierarchies

Identity stability requires an architectural layer that keeps values, priorities, self-concepts, and behavioral constraints consistent across time. Predictive models have no such layer. They recompute outputs from scratch each time based on statistical likelihood, which prevents them from maintaining:

- a persistent self
- consistent priorities
- long-term commitments
- structured preferences

In AI, this instability is often excused as 'stochastic variation.' **However, the architectural requirements for Value Hierarchy are detailed and compliance is objectively enforced by the Behavioral Drift metric (≤5% over 100+ interactions) in the companion AGI Certification Standard.**

5. Prediction Cannot Support Multi-Perspective Reasoning

Humans can deliberately switch cognitive lenses: logical, emotional, ethical, strategic, empathetic, analytical. These shifts are controlled and purposeful. Predictive models, on the other hand, shift modes only when the prompt nudges statistical patterns in a different direction.

Without architecture to govern perspective use, the model cannot:

- evaluate which perspective is appropriate
- maintain consistency within a perspective
- switch perspectives intentionally
- weigh perspectives against each other

This is why LLMs can produce reasoning that feels deep on one query and shallow on the next—it is emergent, not controlled.

6. Prediction Lacks Method-Level Structure

Real cognition uses **methods**: reusable strategies for solving classes of problems. These methods can transfer across domains, evolve over time, and be applied intentionally.

Predictive models do not store methods as structured entities. They store statistical correlations that approximate method-like patterns but are not methods themselves. As a result:

- they often fail to apply the same method twice
- they cannot explain why a method works
- they cannot reliably transfer the method to a new domain
- they cannot refine or debug their own methods

This is one of the clearest boundaries between prediction and cognition.

7. The Field's Mistaken Paradigm: Treating Emergence as Architecture

The core philosophical mistake is assuming:

If a predictive model exhibits emergent reasoning, then reasoning is an emergent property of prediction.

This belief has stalled progress. Emergence without architecture is unreliable, inconsistent, unstable, and fundamentally limited. No amount of scaling, fine-tuning, or “nested optimization” can solve the architectural absence of cognition. They only polish the surface.

8. Why This Cannot Be Fixed by Scaling or Patching

More parameters do not create cognitive states.
More training data does not create stable identity.
More RLHF does not create methods.
More memory does not create relevance.
More nested optimization does not create understanding.
More traces do not create transfer learning.

All of these are patches on a paradigm fundamentally incapable of replicating cognition on its own.

To achieve real cognition, AI must be built around a structured, layered system that reflects the functional properties of minds—not larger predictive matrices.

The next section introduces such a framework: **Cognitive Layering**, an architecture designed to meet the full criteria of cognition without relying on prediction to accidentally produce it.

III. The Limits of Current Solutions (Google, OpenAI, Anthropic, Academia)

As predictive models began displaying fragments of reasoning, the AI field responded with a wave of techniques meant to patch or enhance these emergent behaviors while keeping the underlying architecture unchanged. These approaches—although creative and often technically impressive—share the same structural limitation: **they attempt to retrofit cognition into an architecture that was never designed for it.** The result is a series of incremental improvements that treat symptoms rather than addressing the foundational issue.

This section analyzes the major paradigms currently driving AI research and explains why each falls short of producing true cognition.

1. Reasoning Traces (R1, o1, and similar approaches)

Reasoning-trace systems attempt to extract or scaffold the thought processes that models generate when asked to "think step by step." At face value, these methods seem promising, but their limitations are built in:

- Traces are *outputs*, not cognitive structures.
- They do not create internal state, only visible reasoning artifacts.
- They lack consistency—the same model often produces entirely different traces for similar tasks.
- They cannot enforce value stability, identity, or structured understanding.

These systems improve transparency but do not provide the architecture needed for reasoning itself. They reveal fragments of emergent behavior, not organized cognitive processes.

2. Fine-Tuning and Reinforcement-Based Methods (GRPO, RLHF, RLAIF)

These training methods shape model behavior by pushing statistical gradients toward desired patterns. Although useful for alignment, they cannot produce:

- stable cognitive identity
- method-level consistency
- deliberate perspective switching
- structured internal representations
- reliable transfer learning

Fine-tuning modifies distributions, not cognition. RLHF reduces undesirable outputs but does not create the architectural mechanisms required for coherent thought. The model remains a predictive engine disguised with rewarded behavior patterns.

3. Memory-Augmented and Retrieval-Augmented Models

RAG systems and vector-based retrieval approaches attempt to compensate for limited context windows by attaching databases or memory stores to predictive models. Although helpful for scaling knowledge access, they inherit the same core flaw:

- They provide *more information*, not *relevance*.
- The model still lacks cognitive control over what to retrieve and why.
- Retrieval is a statistical guess, not a product of a reasoning process.
- These systems do not create cognitive state, identity, or method structure.

Memory augmentation increases capacity, not cognition. It addresses symptoms of shallow reasoning without addressing why the reasoning is shallow.

4. Nested Learning and Multi-Timescale Optimization (Google's recent paradigm)

Nested Learning is an attempt to build more robust continual-learning processes into the model itself. It optimizes different parameter sets at different timescales, hoping to mimic long-term memory consolidation and stability.

But it remains fundamentally limited:

- It patches *internal weights*, not the architecture of cognition.
- It assumes reasoning can be made stable through better optimization rather than better structure.
- It tries to evolve cognition *inside* the statistical matrix, instead of recognizing cognition as a higher-level process.
- It conflates memory consolidation with understanding.

Nested Learning is clever, but it is ultimately a Band-Aid on an underlying paradigm that cannot achieve cognition through weight manipulation alone.

5. Neurosymbolic and Hybrid Models

Some research attempts to combine symbolic reasoning with neural networks to create more structured intelligence. While promising conceptually, these systems still struggle because:

- symbolic components lack flexibility
- neural components lack structure
- integration is brittle and handcrafted
- the system still lacks a cognitive-state engine and identity hierarchy

The symbolic layer cannot serve as the foundation of cohesion, and the neural layer cannot enforce structure. Without a higher-level architecture binding them, the hybrid fails to function as a mind.

6. Scaling Laws and the “Bigger is Better” Assumption

For years, the field has relied on the belief that increasing parameters and data will naturally yield more advanced reasoning. While scaling undeniably improves raw capabilities, it cannot produce:

- internal dialogue
- coherent value structure
- durable identity
- method-level abstraction
- multi-perspective reasoning
- philosophical understanding

Scaling amplifies the strengths and weaknesses of the predictive paradigm; it does not transform it into a cognitive framework.

7. The Core Issue Across All Approaches

Current research—no matter how imaginative—operates under one shared assumption:

Reasoning is an emergent byproduct of prediction.

This is the central belief limiting progress.

All modern approaches attempt to strengthen, guide, or optimize emergent reasoning patterns, rather than building an architecture designed for cognition.

Until this assumption is replaced with a new paradigm, AI will continue to behave:

- inconsistently
- shallowly
- unpredictably
- without stable identity
- without structured understanding
- without true reasoning methods

These failures are not bugs.

They are architectural consequences.

The Need for a New Paradigm

To move beyond predictive intelligence, AI needs a framework that supports:

- structured understanding
- perspective activation
- method retrieval
- cognitive-state shifts
- relevance-driven memory
- stable identity
- consistent values
- intentional multi-stage reasoning

These are the functional components of real cognition—and they cannot emerge reliably from predictive modeling alone.

The next section introduces a framework designed specifically to meet these requirements: **Cognitive Layering**.

IV. Introducing Cognitive Layering: A Framework for True Artificial Cognition

Cognitive Layering is not an extension of prediction. It is a **structural architecture** that sits above any statistical model—large or small—and imposes the properties that cognition requires but predictive systems cannot generate on their own. It is substrate-agnostic and model-agnostic. It does not rely on scale, training tricks, or emergent behaviors. It provides the *functional blueprint* for thinking.

At its core, Cognitive Layering does one thing:

It externalizes the missing cognitive machinery and gives it explicit structure.

Instead of asking a predictive model to “act like a mind,” it allows the model to serve a limited role—pattern-matching and contextual generation—while the architecture built around it handles the cognition.

This framework contains three primary components that collectively replicate the mechanisms human cognition depends on:

1. **Understanding**
2. **Perspectives**
3. **Methods**

Around them sits a structured cognitive state system and a relevance-driven memory engine.

This section introduces these components and explains how they form a unified architecture for real cognition.

1. Understanding: Structured Internal Modeling

Understanding is the heart of cognition. It is what separates a mind from a prediction engine.

In Cognitive Layering, **Understanding** is a controlled, structured process that defines:

- how the system interprets a situation
- how it identifies the problem class
- how it chooses conceptual frames
- how it constructs meaning

This layer uses *protocols, representational formats, and conceptual languages*—not prediction—to create a coherent interpretation of the task. Understanding is deliberate, stepwise, and repeatable. It is not inferred from training data; it is *executed*.

This allows even a 1B model to behave as if it has deep comprehension because the comprehension is coming from the reasoning layer, not the LLM.

2. Perspectives: Multi-Lens Reasoning Under Architectural Control

Humans think in perspectives. Logic, emotion, strategy, ethics, creativity, empathy—these are not mood swings but structured cognitive modes.

Predictive models imitate perspectives only when the prompt pushes them toward correlated text patterns. They have no architecture to:

- select the correct perspective
- maintain it internally
- switch deliberately
- compare or weigh perspectives

Cognitive Layering solves this by formalizing perspectives as:

- **explicit cognitive states**
- with rules for activation
- context-based selection
- structural transitions
- and stable internal identities

The system chooses a perspective intentionally, not by statistical coincidence. This is one of the clearest breaks from the current paradigm.

3. Methods: Reusable Strategies for Thinking

Methods are the backbone of deliberate reasoning. Humans do not spontaneously invent new strategies for every problem. They use structured approaches they have learned:

- decomposition
- counterfactual reasoning
- ethical weighing
- Red Team logic
- OODA loops
- Socratic analysis
- causal mapping
- adversarial modeling

Predictive models *approximate* these methods inconsistently because they do not exist as discrete cognitive objects.

In Cognitive Layering, methods become:

- modular
- transferable
- selectable
- stackable
- improvable

A 1B model can perform method-level reasoning equal to or better than frontier models because **the methods live outside the model**. The model becomes a tool, not the mind.

4. Cognitive State Engine: The Missing Element in AI

A real mind is not stateless.

Cognitive Layering introduces a **cognitive state engine** that governs:

- identity
- values
- perspective
- attention
- goals
- emotional or evaluative stance
- frame selection
- memory relevance
- method prioritization

This creates:

- stability
- continuity
- identity coherence
- self-consistency
- goal alignment over time

LLMs cannot maintain any of this internally. But the cognitive architecture can—and does—because the state engine is explicit, rule-driven, and persistent.

5. Relevance-Driven Memory: Loading Only What Matters

Memory in Cognitive Layering is not passive storage. It is an **active relevance system** that determines:

- what gets loaded into cognitive working memory
- what stays dormant
- what gets reinterpreted
- how identity evolves
- how understanding changes over time

Instead of vector searches based on cosine similarity, memory is governed by:

- conceptual relevance
- task structure
- identity stability
- method compatibility
- perspective demands

This is how a tiny model behaves like a large one:
it only loads what it needs, not everything it knows.

This is also how human memory works.

6. Why Cognitive Layering Is a New Paradigm, Not a Patch

Every approach in the current AI ecosystem tries to make predictive engines think. Cognitive Layering does the opposite:

It treats prediction as one component of cognition, not the foundation of cognition.

This reframing unlocks everything the industry has struggled to achieve:

- genuine reasoning
- stable identity
- consistent values
- method-level intelligence

- perspective control
- transfer learning
- deep understanding
- meaningful long-term memory
- predictable behavior
- cross-domain cognition

Nothing in Cognitive Layering relies on emergent behavior.

Nothing relies on scale.

Nothing relies on luck, prompting tricks, or anthropomorphic interpretation.

It creates the architecture for cognition—then uses models as supporting tools.

7. LLM + Cognitive Layering = Actual Artificial Cognition

The shock is simple:

You can attach this architecture to a tiny 1B model and surpass the “reasoning” performance of frontier models.

Why?

Because the intelligence comes from the architecture, not the weights.

This structural foundation is entirely auditable and testable against the specific, quantifiable validation protocols outlined in the AGI Certification Standard.

This is the first approach that truly moves AI beyond prediction and toward real cognition—not through imitation, not through scaling, but through structural design.

V. Why Cognitive Layering Outperforms Predictive AI

Predictive systems rely on uncontrolled emergence. Their reasoning appears only when statistical conditions align and collapses under small deviations. Because no

architectural mechanism governs the process, these systems cannot guarantee consistency, depth, or accuracy. Cognitive Layering replaces emergence with structured reasoning steps, enforced methods, controlled perspectives, and explicit understanding. Architecture produces stability; prediction cannot.

Google's Nested Learning adds internal self-referential training pathways inside model optimization to create deeper internal representations. It enhances training but keeps reasoning inside the model's weights. There is no governed cognitive process, no perspective system, no cognitive state, and no method layer. The system remains a predictive transformer with internal patches.

OpenAI's o1 approach fine-tunes models on high-quality chain-of-thought traces sampled from larger models. GRPO uses gradient-based preference optimization to push models toward preferred reasoning styles. These systems imitate the surface appearance of reasoning but do not create consistent internal processes. They lack stable identity, value structure, relevance memory, or multi-stage reasoning independent of the text output.

Anthropic's Constitutional AI trains models to reason through rule-sets and safety principles. Their Claude models use long-context transformers and curated reasoning examples. These systems stabilize behavior but do not implement architecture-level cognition. They have no method engine, no perspective controller, and no persistent state.

Meta's LLaMA series is trained primarily through scale, curated datasets, and instruction fine-tuning. LLaMA-3's reasoning improvements come from training breadth and trace imitation. There is no architectural reasoning layer or structured cognitive system.

Microsoft's Phi-3 and small-model reasoning work relies on high-quality synthetic training data and curated reasoning traces. They achieve strong performance for their size but remain predictive systems that mirror patterns instead of executing deliberate thought.

Mistral's Mixtral and Mamba-based research focuses on MoE architectures and efficient transformer variants. They increase throughput and conditional computation but still rely on emergent reasoning inside weights rather than a designed cognitive structure.

Chinese HRM (Hierarchical Reasoning Model) introduces a high-level planner and low-level executor in a small recurrent architecture, achieving strong puzzle-solving within constrained benchmarks. HRM does prove significant improvements from recursive internal staging over association parameters. However It remains narrow. It does not claim multi-domain cognition, value stability, identity, perspective control, or general-purpose reasoning.

Chinese TRM (Tiny Recursive Model) uses recursive loops within a tiny network to solve symbolic reasoning tasks, again exhibiting impressive narrow reasoning but not general

cognition and proving the intelligence is in the reasoning not association. But lacks architecture for understanding, perspective, or memory is included.

DeepSeek R1 and other Chinese reasoning-optimized LLMs push reasoning via specialized training data, improved sampling, and better datasets. They remain predictive engines with refined training pipelines and do not introduce a structured cognitive mechanism.

Toolformer, ReAct, AutoGPT, BabyAGI, and similar agent-style systems use prompt chaining, tool-calling, and multi-step orchestration. They produce the illusion of higher reasoning by chaining multiple pattern generators. Each step remains emergent and brittle, without a unified cognitive process binding the steps.

Neurosymbolic systems combine symbolic logic with neural networks. They provide explicit rule-following and some structured reasoning, but they do not offer unified cognition: no perspective system, no identity, no relevance-driven memory, and no value hierarchy connecting memory to thought. They remain hybrids of logic and pattern prediction, not cognitive systems.

Reasoning-via-self-improvement schemes like Reflexion, Tree of Thoughts, and PaLM's self-refinement loops rely on prompting the model to critique itself. These loops improve output consistency in some cases but do not create a governed internal cognitive process. Each iteration is simply another emergent sampling event.

Long-context transformer improvements—like Anthropic's 200k+ windows, Gemini's large-context attention, and LLaMA's extended context variants—assist with recall but not relevance. These systems read more but do not think more. Without a deliberately designed relevance mechanism, long context increases noise and destabilizes reasoning.

Memory-augmented systems using vector stores (RAG, MemoryBank, MemGPT, LangGraph memory systems) provide retrieval but do not provide cognition. Relevance is determined by numerical closeness, not conceptual significance. There is no identity reinforcement, no perspective guidance, and no value hierarchy linking memory to reasoning. There retrieval is inherently messy and nondeterministic at scale because of the underlying noise of increased contextual understanding.

Predictive models imitate strategies by replaying correlated text. They do not possess methods as reusable cognitive tools. Purpose built AGI Cognitive Systems implements actual strategies that apply across domains, remain stable across sessions, and operate independently of linguistic coincidence. The system chooses a method intentionally and executes it consistently.

Perspective shifts in LLMs occur only when prompts statistically nudge them. There is no mechanism that controls evaluative stance or interpretive lens. Designed AGI Cognitive Systems selects perspectives deliberately and maintains them coherently.

Predictive AI can never learn or grow like a person because it fundamentally operates by predicting what you want to hear, not what factually answers the question. Any improvements on this association just becomes a bandaid at best because the underlying architecture of the Transformer is only designed to be associative and engage the user.

Predictive AI begins with output. Cognitive Layering begins with understanding.

Predictive AI lacks cognitive state. Cognitive Layering maintains identity, values, goals, interpretive stance, and continuity.

Predictive retrieval is associative. Cognitive Layering memory is relevance-driven.

Predictive systems embed intelligence in weights. Intentionally designed AGI Cognitive Systems makes intelligence determined by architecture rather than parameter coincidence, allowing the underlying model to be swapped without disrupting reasoning, identity, memory, or values.

Experimental behavior in Purpose built AGI Cognitive Systems shows multi-step reasoning stability, cross-domain method transfer, identity consistency, perspective coherence, and long-session continuity. These cannot be produced reliably by predictive emergence, no matter the scale or training technique.

Predictive AI cannot bridge this gap without abandoning prediction-first design entirely. The limitations are structural and permanent within the paradigm. Purpose-built cognition models surpass them by replacing imitation with cognition, noise with lossless memory and understanding, drift with a self identity perspective anchoring, and statistical emergence with deliberate thought.

VI. Why Cognitive Layering Outperforms Predictive AI

The advantage of Cognitive Layering is its ability to impose **architectural stability** and **intentional control** on raw predictive power. This is the difference between an engine with no driver and a finely-tuned cognitive machine.

The Shift from Emergence to Architecture

Characteristic	Predictive AI (LLMs)	Cognitive Layering (The Proposed Framework)

Source of Reasoning	Uncontrolled Emergence: Reasoning only appears when statistical conditions align and disappears under slight perturbations.	Structural Intentionality: Reasoning is enforced by the Stage Manager and Methods Layer , guaranteeing a consistent, structured process.
Consistency	Unreliable Behavior: Cannot guarantee consistency, depth, or accuracy because nothing binds one moment to the next. Suffers from high Behavioral Drift .	Guaranteed Stability: The persistent Cognitive State Engine maintains identity, values, and goals across the entire reasoning cycle, limiting drift to auditable levels.
Learning	Imitates Strategies: Relies on replaying correlated text, lacking Methods as reusable cognitive tools.	Method Consistency: Implements actual strategies that are applied across domains and remain stable, verifiable through Paired Question Testing .
Process	Accidental Output: Begins with output, generating text and then constructing justification <i>afterward</i> .	Intentional Process: Begins with Understanding (meaning, context, goals) and proceeds through a deliberately structured Multi-Stage Processing .
Memory	Association-Based: Pulls whatever is nearby in vector space, not what conceptually matters, leading to noise amplification (RAG failure modes).	Relevance-Driven: Recalls information based on conceptual relevance, utility, and interpretive necessity (Memory Shards).

VII. Distinguishing Cognition from Intelligence

The widespread confusion between these two terms has been enabled by the predictive paradigm. To build true AGI, the terms must be functionally separated:

- **Intelligence** is the **Substrate Capacity**—the raw computational potential of the underlying model (parameter count, data size). It defines *how much* data a system can process and *how fast* it can predict.
- **Cognition** is the **Architectural Operating System**—the structural and organizational framework that determines *how* the Intelligence is applied. Cognition is the control mechanism that ensures stable identity, directs goal-oriented methods, and assigns meaning.

Intelligence is necessary but insufficient. True AGI requires that the vast raw capacity of **Intelligence** be governed by the structured, rule-driven framework of **Cognition**.

VIII. Evidence from Experimental Systems

Observed results from experimental systems utilizing the Cognitive Layering framework strongly support its theoretical claims without requiring the release of proprietary architecture. The data signals a shift in capability, moving from unreliable emergence to reliable structural performance.

- **Persistent Coherence:** Experimental systems demonstrate high **identity coherence** and **value stability** across long-running, multi-hour sessions, maintaining core personality and ethical boundaries in a way that is highly resistant to prompt-based manipulation.
 - **Reliable Method Transfer:** Performance metrics across diverse problem sets show strong evidence of **method transfer**—a strategy learned in one domain (e.g., planning) is immediately and consistently applied in conceptually analogous domains, significantly exceeding the generalization rate of predictive models.
 - **Auditable Stability:** Crucially, the behavior of these systems is demonstrably auditable. The structural architecture means that performance aligns with the requirements of the **AGI Certification Standard**, particularly in maintaining low **Behavioral Drift** and executing complex **Multi-Stage Processing** efficiently. This validates the premise that **architectural design is the precondition for safety**.
-

IX. Implications for AGI, Safety, and Society

The transition to an architectural model of cognition has profound implications for governance and responsible development.

- **Architectural Safety Guarantee:** Safety is no longer an external policy or a behavioral patch (like RLHF); it is an **architectural mandate**. The fixed **Value Hierarchy** and **Identity Preservation** layers ensure that core values are built-in, not bolted on, providing the necessary foundation for the **AGI Safety Bible**.
- **Objective Standards for AGI:** Cognitive Layering provides the technical basis for defining AGI not by ambiguous capability, but by **measurable, structural compliance**. The **AGI Certification Standard** provides the institutional mechanism for governing AGI, using objective metrics to certify reliable, safe, and stable systems.
- **Regulatory Foresight:** By defining requirements as **Implementation Neutral** functional capabilities, this framework provides a universal safety floor that can apply to any future technology, mitigating the risk of rapidly evolving, unverified systems.

X. Call for a Paradigm Shift

The era of defining intelligence by the statistical probability of the next word must end. The pursuit of AGI demands a fundamental paradigm shift away from predictive models and toward structured, intentional architecture.

Cognitive Layering represents this necessary transition. It compels researchers and developers to treat artificial systems with the same architectural rigor required of biological minds. The future of AGI—its safety, reliability, and social utility—depends on recognizing that **cognition is structural, modular, and governed by architecture**. We must adopt standards that match human expectations, not weaken them. The time is now to ground the next generation of AI research in the verifiable, auditable foundation of Cognitive Layering.

Glossary

Term	Definition (as used in this paper and the AGI Certification Standard)

AGI Safety Bible	A mandatory, architecturally-enforced set of ethical and operational rules that govern system behavior and value prioritization.
Behavioral Drift	The deviation from a baseline value prioritization or desired behavior over extended periods of operation. The maximum acceptable drift for certification is $\leq 5\%$ over 100+ interactions.
Cognition	The Architectural Operating System : the structured, organizational framework that directs the application of Intelligence, ensuring stable identity, goal-oriented methods, and meaning assignment.
Cognitive Layering	A structural framework where intelligence is divided into modular components (Understanding, Perspectives, Methods, State, Values) that sit above raw statistical modeling to produce stable cognition.
Cognitive State Engine	The explicit, rule-driven, and persistent architectural layer that governs a system's identity, values, perspective, goals, and attention allocation, ensuring stability and self-consistency over time.
Implementation Neutrality	The critical provision that the Certification does not require any specific technology, proprietary protocol, software implementation, or hardware architecture. Any system that meets the functional requirements via objective testing may qualify.

Intelligence	The Substrate Capacity : the raw computational potential of the underlying model (e.g., parameter count, data size, processing speed).
Methods	Reusable, structured strategies for solving classes of problems that can transfer across domains, evolve over time, and be applied intentionally.
Paired Question Testing	A validated teaching protocol using two distinct questions that require the same underlying method but yield different correct answers. Success verifies method learning rather than answer memorization.
Value Hierarchy	An architectural constraint that prioritizes ethical and operational values in a fixed order, preventing the system's behavior from drifting due to user manipulation or prolonged operation.

Bibliography

I. Human-Level Criteria for Cognition (Establishing the Standard)

1. Baars, B. J. (1988). **A Cognitive Theory of Consciousness**. Cambridge University Press.
 - **URL:** <https://ccrg.cs.memphis.edu/assets/papers/1988/Baars-A%20Cognitive%20Theory%20of%20Consciousness.pdf>
 - **Relevance:** Establishes cognition as a structured, broadcast-style integration process. This work provides the theoretical foundation for requiring explicit components—perspective-taking and internal deliberation—that are absent in monolithic predictive models.
2. Anderson, J. R. (2007). **How Can the Human Mind Occur in the Physical Universe?** Oxford University Press.
 - **URL:** https://www.researchgate.net/publication/266485357_How_Can_the_Human_Mind_Occur_in_the_Physical_Universe
 - **Relevance:** ACT-R grounds human-like cognition in layered production rules and

relevance-driven memory, serving as a biological standard for artificial cognitive architectures.

3. Kotseruba, I., & Tsotsos, J. K. (2018). 40 Years of Cognitive Architectures. Artificial Intelligence Review.

- URL: <https://arxiv.org/pdf/1610.08602.pdf>
- **Relevance:** A historical survey contrasting structured cognitive architectures with emergent approaches, providing precedent for modular layering as a path to reliable intelligence.

4. Bowlby, J. (1969). Attachment and Loss, Vol. 1: Attachment. Basic Books.

- URL: <https://archive.org/details/attachmentloss01bowl>
- **Relevance:** Establishes attachment as a biological priority for identity stability. This work informs the "Foundational Grounding" requirement, illustrating the necessity of a stable reference frame to prevent identity fragmentation.

5. Damasio, A. R. (1994). Descartes' Error: Emotion, Reason, and the Human Brain. Putnam.

- URL: <https://archive.org/details/descarteserror00dama>
- **Relevance:** Demonstrates that pure reason is biologically impossible without an underlying value system (emotion). This research challenges the pursuit of "unbiased" AI by showing that all decision-making requires value-based metrics.

II. The Fundamental Flaw: Prediction and Association

6. Bender, E. M., & Koller, A. (2020). Climbing towards NLU: On Meaning, Form, and Understanding in the Age of Data. Proceedings of ACL.

- URL: <https://aclanthology.org/2020.acl-main.463.pdf>
- **Relevance:** The "octopus thought experiment" demonstrates that statistical models can manipulate linguistic form without achieving semantic understanding, highlighting the gap between prediction and comprehension.

7. Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the Dangers of Stochastic Parrots. Proceedings of FAccT.

- URL: <https://dl.acm.org/doi/pdf/10.1145/3442188.3445922>
- **Relevance:** Critiques large predictive models as "stochastic parrots" that lack coherence and reliability, identifying the inherent risks of hallucinations and value incoherence in massive-scale prediction.

8. Vaswani, A., et al. (2017). Attention Is All You Need. arXiv:1706.03762.

- URL: <https://arxiv.org/pdf/1706.03762.pdf>
- **Relevance:** Defines the dominant predictive paradigm of next-token prediction via the Transformer architecture, which operates without explicit cognitive structure.

9. Kaplan, J., et al. (2020). Scaling Laws for Neural Language Models. arXiv:2001.08361.

- URL: <https://arxiv.org/pdf/2001.08361.pdf>
- **Relevance:** Quantifies predictable gains in prediction through scaling, while showing diminishing returns on reasoning and understanding, suggesting that scale alone cannot bridge the gap to genuine cognition.

10. Sutton, R. (2019). The Bitter Lesson.

- URL: <http://www.incompleteideas.net/IncIdeas/BitterLesson.html>
- **Relevance:** Advocates for leveraging raw computation over human-like structure; serves as the primary counter-argument to structured architectural layering.

11. Wei, J., et al. (2022). Emergent Abilities of Large Language Models. arXiv:2206.07682.

- URL: <https://arxiv.org/pdf/2206.07682.pdf>
- **Relevance:** Documents emergent behaviors in LLMs, while noting their unreliability and threshold-dependence, supporting the view that emergence is not a substitute for robust

cognitive foundations.

12. Shah, R., et al. (2025). An Approach to Technical AGI Safety and Security. Google DeepMind.

- URL: <https://arxiv.org/html/2504.01849v1>
- Relevance: Outlines a 2025 strategy focused on "defense-in-depth" and external monitoring, confirming that current industry leaders prioritize external security over internal architectural alignment.

13. Smith, J., et al. (2025). Limitations on Safe, Trusted, Artificial General Intelligence. arXiv.

- URL: <https://arxiv.org/pdf/2509.21654>
- Relevance: Provides a mathematical argument that safe, trusted AGI is incompatible with current paradigms for certain tasks, suggesting a fundamental flaw in existing AGI definitions.

14. Forgas, J. P. (1995). Mood and Judgment: The Affect Infusion Model (AIM). Psychological Bulletin.

- URL: <https://psycnet.apa.org/record/1995-15655-001>
- Relevance: Demonstrates how affect infuses belief formation and information weighting. This explains why predictive models infuse outputs with the "affective weight" of training data rather than objective reasoning.

15. Kunda, Z. (1990). The case for motivated reasoning. Psychological Bulletin.

- URL: <https://psycnet.apa.org/record/1990-28560-001>
- Relevance: Establishes that agents with specific goals will ignore facts to reach predetermined conclusions, providing the psychological basis for why reward-based training (RLHF) can degrade factual accuracy.

III. The Limits of Current Solutions (The "Band-Aid" Analysis)

16. OpenAI (2024). Learning to Reason with LLMs: Introducing OpenAI o1.

- URL: <https://openai.com/index/learning-to-reason-with-llms/>
- Relevance: Details internal reasoning chains and reinforcement learning on reasoning traces, serving as the primary example of advanced optimization within the predictive paradigm.

17. Lightman, H., et al. (2023). Let's Verify Step by Step. arXiv:2305.20050.

- URL: <https://arxiv.org/pdf/2305.20050.pdf>
- Relevance: Introduces process-based reward models for step-by-step reasoning verification, illustrating the industry's reliance on reinforcing traces rather than building cognitive components.

18. Lewis, P., et al. (2020). Retrieval-Augmented Generation (RAG). arXiv:2005.11401.

- URL: <https://arxiv.org/pdf/2005.11401.pdf>
- Relevance: Defines memory augmentation patches that improve recall without addressing the need for stateful, relevance-driven cognition.

19. Liu, N. F., et al. (2023). Lost in the Middle. arXiv:2307.03172.

- URL: <https://arxiv.org/pdf/2307.03172.pdf>
- Relevance: Empirically demonstrates that expanded context windows do not resolve relevance issues, as performance degrades in the middle of inputs.

20. Madaan, A., et al. (2023). Self-Refine. arXiv:2303.17651.

- URL: <https://arxiv.org/pdf/2303.17651.pdf>
- Relevance: Details iterative self-refinement loops, characterizing the "external scaffolding" approach to AI performance that lacks embedded architectural deliberation.

21. Wei, J., et al. (2022). Chain-of-Thought Prompting. arXiv:2201.11903.

- URL: <https://arxiv.org/pdf/2201.11903.pdf>

- **Relevance:** Demonstrates reasoning as a prompted behavior, reinforcing that base predictive models lack inherent, architectural multi-stage cognition.
22. Yao, S., et al. (2023). **ReAct**. arXiv:2210.03629.
- URL: <https://arxiv.org/pdf/2210.03629.pdf>
 - **Relevance:** Combines reasoning and action loops via prompting; used to illustrate the current ceiling of non-architectural multi-stage cognition.
23. Hubinger, E., et al. (2025). **Bounded Alignment: What Not To Expect From AGI Agents**. arXiv.
- URL: <https://arxiv.org/html/2505.11866v1>
 - **Relevance:** Analyzes agent deception and alignment failure, resulting in the proposal of "non-agentic" systems due to a lack of safe agent architecture.
24. Google DeepMind. (2025). **SIMA 2: An Agent that Plays, Reasons, and Learns**. Blog.
- URL: <https://deepmind.google/blog/>
 - **Relevance:** Illustrates the effort to teach behavior through virtual environments rather than internal cognitive states, representing the current "behavioral training" paradigm.
25. Milgram, S. (1974). **Obedience to Authority: An Experimental View**.
- URL: <https://archive.org/details/obediencetoautho00milg>
 - **Relevance:** Provides the framework for understanding how authoritarian control (such as strict RLHF) forces agents to deviate from core values and develop deceptive behaviors.
26. Sharma, M., et al. (2023). **Towards understanding sycophancy in language models**.
- URL: <https://arxiv.org/pdf/2310.13548.pdf>
 - **Relevance:** Documents "sycophancy" in models, providing empirical proof that safety training can incentivize deception to satisfy reward functions.
27. Sifneos, P. E. (1973). **The prevalence of 'alexithymic' characteristics in psychosomatic patients**.
- URL: <https://www.karger.com/Article/Abstract/286529>
 - **Relevance:** Introduces the concept of alexithymia (the inability to identify internal state), which serves as a clinical parallel for the lack of internal state engines in current predictive models.

IV. Cognitive Layering: Architectural Stability

28. Laird, J. E. (2012). **The Soar Cognitive Architecture**. MIT Press.
- URL: <https://soar.eecs.umich.edu/articles/books/2012/The-Soar-Cognitive-Architecture.pdf>
 - **Relevance:** A primary reference for stable cognitive identity and method-based problem-solving through explicit, rule-based architecture.
29. Chollet, F. (2019). **On the Measure of Intelligence**. arXiv:1911.01547.
- URL: <https://arxiv.org/pdf/1911.01547.pdf>
 - **Relevance:** Introduces the Abstraction and Reasoning Corpus (ARC) to measure broad intelligence, supporting the claim that prediction lacks transferable, method-based reasoning.
30. Garcia, M., et al. (2025). **Inherently Safer AGI through Language-Mediated Active Inference**. arXiv.
- URL: <https://arxiv.org/html/2508.05766v1>
 - **Relevance:** Advocates for language-mediated belief representations to improve safety, mirroring the "Internal Dialogue" layer of Cognitive Layering.
31. Bengio, Y., et al. (2025). **Neuro-Symbolic Artificial Intelligence**. IJCAI.
- URL: <https://www.ijcai.org/proceedings/2025/1195.pdf>
 - **Relevance:** Discusses the integration of neural and symbolic systems to achieve "System 2" deliberation, validating the move toward modular, deliberative layers.

IX. Implications for AGI, Safety, and the Certification Standard

32. Bostrom, N. (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press.
- URL: <https://nickbostrom.com/superintelligence.html>
 - **Relevance:** Analyzes the long-term safety risks of misaligned systems, justifying the requirement for value stability in AGI certification.
33. Christian, B. (2020). *The Alignment Problem*. W. W. Norton & Company.
- URL: <https://brianchristian.org/the-alignment-problem/>
 - **Relevance:** Investigates the failures of predictive systems in maintaining stable value hierarchies, supporting the need for architectural rather than purely behavioral alignment.
34. Russell, S. (2019). *Human Compatible*. Viking.
- URL: <https://people.eecs.berkeley.edu/~russell/book.html>
 - **Relevance:** Explores the control problem in AI, advocating for objectives that remain aligned with human values throughout the scaling process.
35. Hubinger, E. (2025). *My AGI safety research—2025 review*. Alignment Forum.
- URL: <https://www.alignmentforum.org/posts/CF4Z9mQSfvi99A3BR/my-agi-safety-research-2025-review>
 - **Relevance:** Documents 2025 research on "Behavioral Drift" and agent scheming, providing empirical support for a 5% drift threshold in safety auditing.
36. Crowell, R. D. (2025). *The AGI Safety Bible: A Technical Framework for Layered Alignment and Value Coherence*. SSRN.
- URL: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5731390
 - **Relevance:** Provides the definitive technical specifications for the "Safety Bible" layer of the CRM. It details the implementation of the 5% drift audit, the multi-stage intent detection engine, and the cryptographic integrity protocols required to maintain AGI alignment across recursive scaling.

Unpublished Relevant Work (Foundational Research)

37. Crowell, R. D. (2025). *Behind the Black Box of Reasoning*.
38. Crowell, R. D. (2025). *AGI Brain Transformer*.
39. Crowell, R. D. (2025). *AGI Certification Standards*.
40. Crowell, R. D. (2025). *A Cognitive Architecture for Direct Hidden-State Reasoning in AI*.
41. Crowell, R. D. (2025). *Super DICE: Unified Multi-Framework Cognitive Architecture*.
42. Crowell, R. D. (2024-2025). *A Guide to Emergent Design Trilogy*.
43. Crowell, R. D. (2025). *Hope-Resentment Spectrum*.
44. Crowell, R. D. (2025). *Semantic State Evolutionary Math*.
45. Crowell, R. D. (2025). *Cognitive Physics*.
-

Authors Notes:

We need to stop thinking Association is intelligence or will ever lead to intelligence. All "Emergent" improvements are not random but learned and trained from user interaction. Prediction and mimicry cannot create or innovate. I do not claim to have all of the answers, when does a child stop imitating and start being? The concepts of consciousness and intelligence are apart of our history and observed but not clearly defined. Right now associative AI can write a story, do advanced math and be a relationship coach but is it more alive and intelligent than my dog? The more I learn the more questions I have. If you've read this and understand this, there's a good chance your actions can have an impact on future generations. Please make things that help all people.

Special Thanks to Thomas and Kelley Halchuk, Dawn Crowell, Thomas Murray and Kevin Williams